

Knowledge engineering on PhD stories: bottom-up vs. goal-directed approach

VIET BACH NGUYEN¹, STANISLAV KRUML¹, VOJTĚCH SVÁTEK¹ and ÓSCAR CORCHO²

¹*Prague University of Economics and Business, Nám. W. Churchilla 4, 130 67 Praha 3, Czech Republic
E-mail: viet.nguyen@vse.cz krus02@vse.cz svatek@vse.cz*

²*Universidad Politécnica de Madrid, Madrid, Spain
E-mail: ocorcho@fi.upm.es*

Abstract

Stories of concrete PhD students, reflected in various information sources, are potentially a useful source of reflections and analytics on the underlying generic processes and regulations, and can subsequently give rise to guidelines and hints making the life of future PhD students easier. The collected information on PhD stories is of quite heterogeneous nature, from structured ‘master data’ from study information systems, through various partial records in broader (e.g., publication) databases, to textually expressed chronologies, impressions, and lessons learned prepared by the PhD students themselves. We explored two fundamentally different approaches to arrive at a structured representation of individual PhD stories. The first approach is bottom-up, relying on aspect-based sentiment analysis (SA) as an established NLP method. Training data was collected in the form of 35 textual stories (or, more broadly, documents reflecting the PhD experience of the author), manually labeled by volunteers, and submitted to state-of-the-art SA tools. The second approach is primarily manual in its nature, though amenable to knowledge/database support: population of a knowledge graph (or, concept map) from primarily structured sources as well as insider information, guided by a hierarchical system of (PhD study) goals. The task was accomplished, as a mere proof of concept, by two senior academics with ample experience in supervising PhD students. Both approaches appear complementary, but insufficient overall, to date. While the bottom-up one can much better scale, the result is largely incomplete in its coverage. The goal-directed approach is, in turn, extremely demanding on human expertise and laborious due to lack of proper tooling. Based on the gained experience, we discuss conditions that could allow making knowledge engineering on PhD stories more efficient and effective in the future.

1 Introduction

The research follows up with an initial pilot Nguyen et al. (2021), which however only featured the clean bottom-up approach in the form of thematic codes being provisionally assigned to twelve documents by a single researcher, and several ad hoc variants of graph modeling, without the support of any explicit guidelines. The current research, in contrast, newly introduces aspect-based sentiment analysis as a particular computational technique, in the bottom-up approach, and standardizes the graph modeling (goal-driven) approach through common guidelines. The total size of the collection has doubled (tripled). Explicit formulation of requirements allowing for effective knowledge engineering in the given domain is yet another contribution of the paper.

The structure of this paper is as follows:

2 Related research

2.1 PhD story analysis

2.2 Academic knowledge graphs

2.3 Biography modeling

3 Bottom-up analysis

3.1 Data collection

The PhD story data used in this research are mostly acquired from websites where they are published by their respective authors or by the permitted publishers. Apart from that, we also have gathered several “offline” stories via face-to-face interviews. Fig. 1 briefly shows our web search strategy along with the proportion of the resources by each document type.

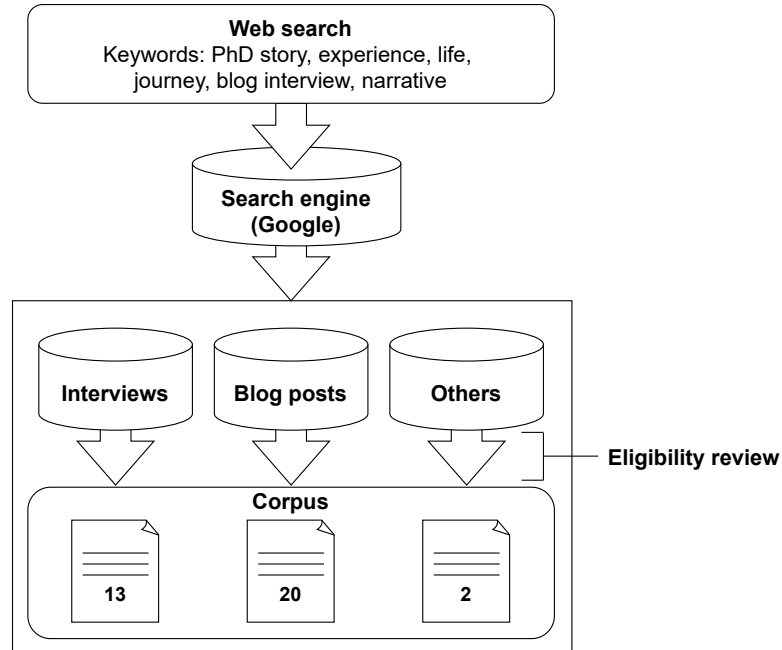


Figure 1 Data collection strategy

To collect stories that are published online, we follow a simple keyword-based search strategy on Google to produce results in the form of links. In this search strategy, suitable keywords are defined to search for PhD stories. These keywords are *PhD story*, *my PhD story*, *my PhD journey*, *PhD narrative*, *PhD experience*, *PhD blog*, *PhD interview*, *PhD life*, and some combinations. Next, we visit each link and quickly scan the story content to determine whether a story is eligible for further analysis. Our selection criteria are not strictly defined, since every story is unique in writing style, form, and thoughts being conveyed. However, the most important first-hand properties we look for in the stories during the web search are language (only English), summarizing intent, sufficient story length (e.g., approximately five hundred words or more), high coverage of multiple PhD life aspects¹, high density in retrospective narrations (e.g., lesson learned, experience, situations, decisions), chronological composition (e.g., sectioned into academic years), and recentness of the story (e.g., we only consider modern PhD experiences, thus only stories published in the last 20 years are collected, this should exclude PhD programs active in the 1990s and earlier).

¹...list the aspects from the aspect analysis phase

In other words, the search results are filtered², and if the data do not satisfy the above-mentioned criteria, they are omitted. In the case of “offline” interviews, we guided the interviewees with our open interview questions to allow them to speak freely but also to gather the needed information. These interview records were later transcribed into text and included in the final dataset. Our dataset is a corpus of a total of 35 eligible PhD stories in textual form with about a total of 76 088 words, ranging from 500 words to 6 000 words each and the average word count is 2 174. We classify the document types into blog posts (20), interviews (13), and others (1 vlog and 1 web article).

3.2 Document categorization

Since the nature of the documents varied, we manually skimmed them so as to establish categories (clusters) of documents having similar natures and/or destinations. ...

3.3 Aspect system design

We design the aspects via brainstorming where we list out different areas and topics that a PhD student may talk about. Also, with the consultancy of our senior member, we have created the following tree of aspects to be used during aspect annotation

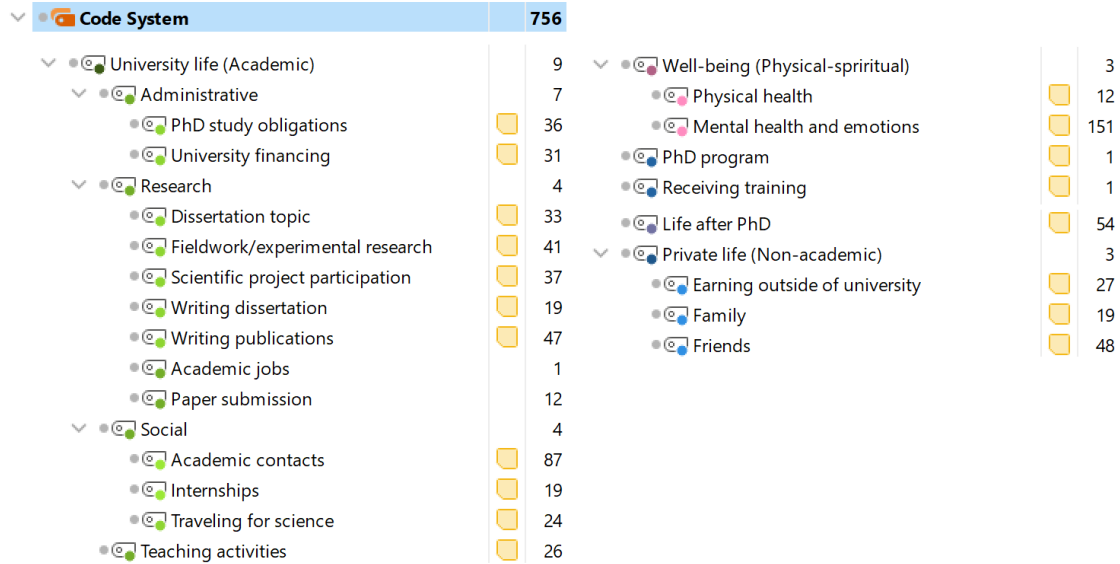


Figure 2 Annotation codes – aspects

3.4 Aspect annotation

We annotate the collection of PhD stories using MAXQDA software. For this annotation of aspects, we had a team of 6 people annotating 24 stories (at the time we only had 24 stories), see Tab. 1. Each person took care of 6 stories so that 12 stories are annotated once and the 12 stories are annotated twice by two people in hopes of measuring the annotators’ agreement. We split the 24 stories into 6 pools evenly so that the time spent in each story pool would be approximately

²We do not filter the stories based on the research field of the authors, but we initially wanted to limit our search scope to Applied Informatics only. We decided to drop this criterion since the initial number of search results was not enough. We then expanded our search scope to the Computer Science domain, but again the search yielded too few results. Therefore, after removing this limitation, our final dataset contains stories from different research domains, such as biology, physics, maths, engineering, literature, etc.

the same as in other pools. The time measured is based on the total word count of the stories in each pool. We approximated the reading and annotating speed to 100 words per minute while the average reading speed of an adult person is around 238 words per minute. The reported time spent in each pool by the annotators turned out to be approximately the same. In this aspect annotation phase, we have produced a total of 756 annotations.

pool	files to annotate	time estimate [m]	word count	annotator
A	09, 13, 17, 18, 20, 21	119	12 182	Vojta
B	02, 06, 08, 13, 14, 19	114	11 630	Standa
C	01, 03, 06, 08, 11, 15	120	12 245	David
D	01, 04, 05, 10, 12, 23	110	11 083	Marek
E	03, 16, 17, 19, 22, 23	123	12 515	Bach
F	07, 10, 12, 15, 21, 24	122	12 419	Gollam

Table 1 Aspect annotation pools

+

3.5 Polarity and emotion annotation

Aside from three-valued polarity, we also used this opportunity to acquire information on the *surprise* emotion. The reason is that the surprise emotion is one that does not inherently lean towards positive nor negative polarity, yet can be an important ingredient in a PhD student’s story – whether in the context of the actual research (succeeding or failing against odds) or, e.g., social relationships within the lab or school.

Code System	1878
mixed	218
negative	545
positive	1060
surprised	55

Figure 3 Annotation codes – sentiments

We also annotate our collection of PhD stories with sentiments using MAXQDA software. This time we had a different team of 8 people manually annotate all 35 stories using 4 annotation codes which are *positive*, *neutral*, *negative*, and the mentioned *surprise* emotion. For this annotation task, we divide our corpus of stories into 8 distinct pools with 6–7 stories each so that 18 stories are annotated twice by two people. Each pool represents about 155 minutes of reading and annotating time for approximately 15 500 words (we stick with the speed of 100 words per minute from the aspect annotation phase). We calculate these estimated time values based on the word count of each document and distribute the stories into the pool evenly by time, see Tab. 2.

pool	files to annotate	time estimate [m]	total words	annotator
A	17, 20, 22, 25, 29, 31, 32	151	15 105	Dominik
B	08, 13, 14, 15, 28, 34, 35	146	14 642	Šimon
C	06, 08, 10, 19, 27, 30	160	16 045	Michael
D	01, 02, 04, 05, 23, 25, 27	161	16 080	Dominika
E	01, 03, 13, 17, 21, 23, 26	149	14 863	Karel
F	07, 10, 12, 15, 24, 33, 35	151	15 069	Maxim
G	03, 06, 09, 18, 19, 29	151	15 148	Kateřina
H	11, 12, 16, 21, 31, 33	159	15 876	Marek

Table 2 Sentiment annotation pools

Annotators were then randomly assigned to these pools. We also provide the annotators with a common guide document explaining the necessary steps for the annotations. Our definitions of sentiment values are the following:

- *positive* – there is an explicit or implicit clue in the sentence suggesting the writer’s attitude towards or judgment of the subject is positive (thankful, excited, optimistic, overjoyed, happy, inspired, etc.),
- *negative* – there is an explicit or implicit clue in the sentence suggesting the writer’s attitude towards or judgment of the subject is negative (critical, angry, disappointed, pessimistic, sarcastic, mocking, bored, complaint, etc.),
- *mixed* – there is an explicit or implicit clue in the sentence suggesting that the writer’s attitude towards or judgment of the subject is both positive and negative,
- *unknown* – there is no explicit or implicit clue indicating that the writer feels positively or negatively about the subject (this value is not used during the annotation process, whatever text that remains after annotating is considered unknown)

Each annotated sentence can only have 1 sentiment (positive, mixed, or negative – these are disjoint). However, these sentences may also include a surprise emotion, such as when there is a sentence containing ”I didn’t know“, or ”it was nothing like I imagined“, etc. In that case, they should also be annotated with the *surprised* emotion code. In this sentiment annotation phase, a total of 1878 annotations are created.

3.6 Sentiment analysis: employed techniques

3.7 Sentiment analysis: results

4 Goal-directed analysis

4.1 Data sources

4.2 Modeling guidelines

4.3 Produced concept graphs

accessibility of Semantic Scholar in ContextMinds

5 Discussion

6 Conclusion

Future: fetching data from DBLP and similar databases

Acknowledgements

This research was supported by VSE IGS F4/56/2021. The authors are indebted to the volunteer annotators, without whom the research described in Section 3 would not have been possible.

References

- Nguyen, V. B., V. Svátek, M. Dudás, and Ó. Corcho
 2021. Knowledge engineering of PhD stories: A preliminary study. In *K-CAP ’21: Knowledge Capture Conference, Virtual Event, USA, December 2-3, 2021*, A. L. Gentile and R. Gonçalves, eds., Pp. 281–284. ACM.